The Quest for Unification:

A Survey of Hardware-Agnostic Machine Learning Systems

Paul Dutton

TABLE OF CONTENTS

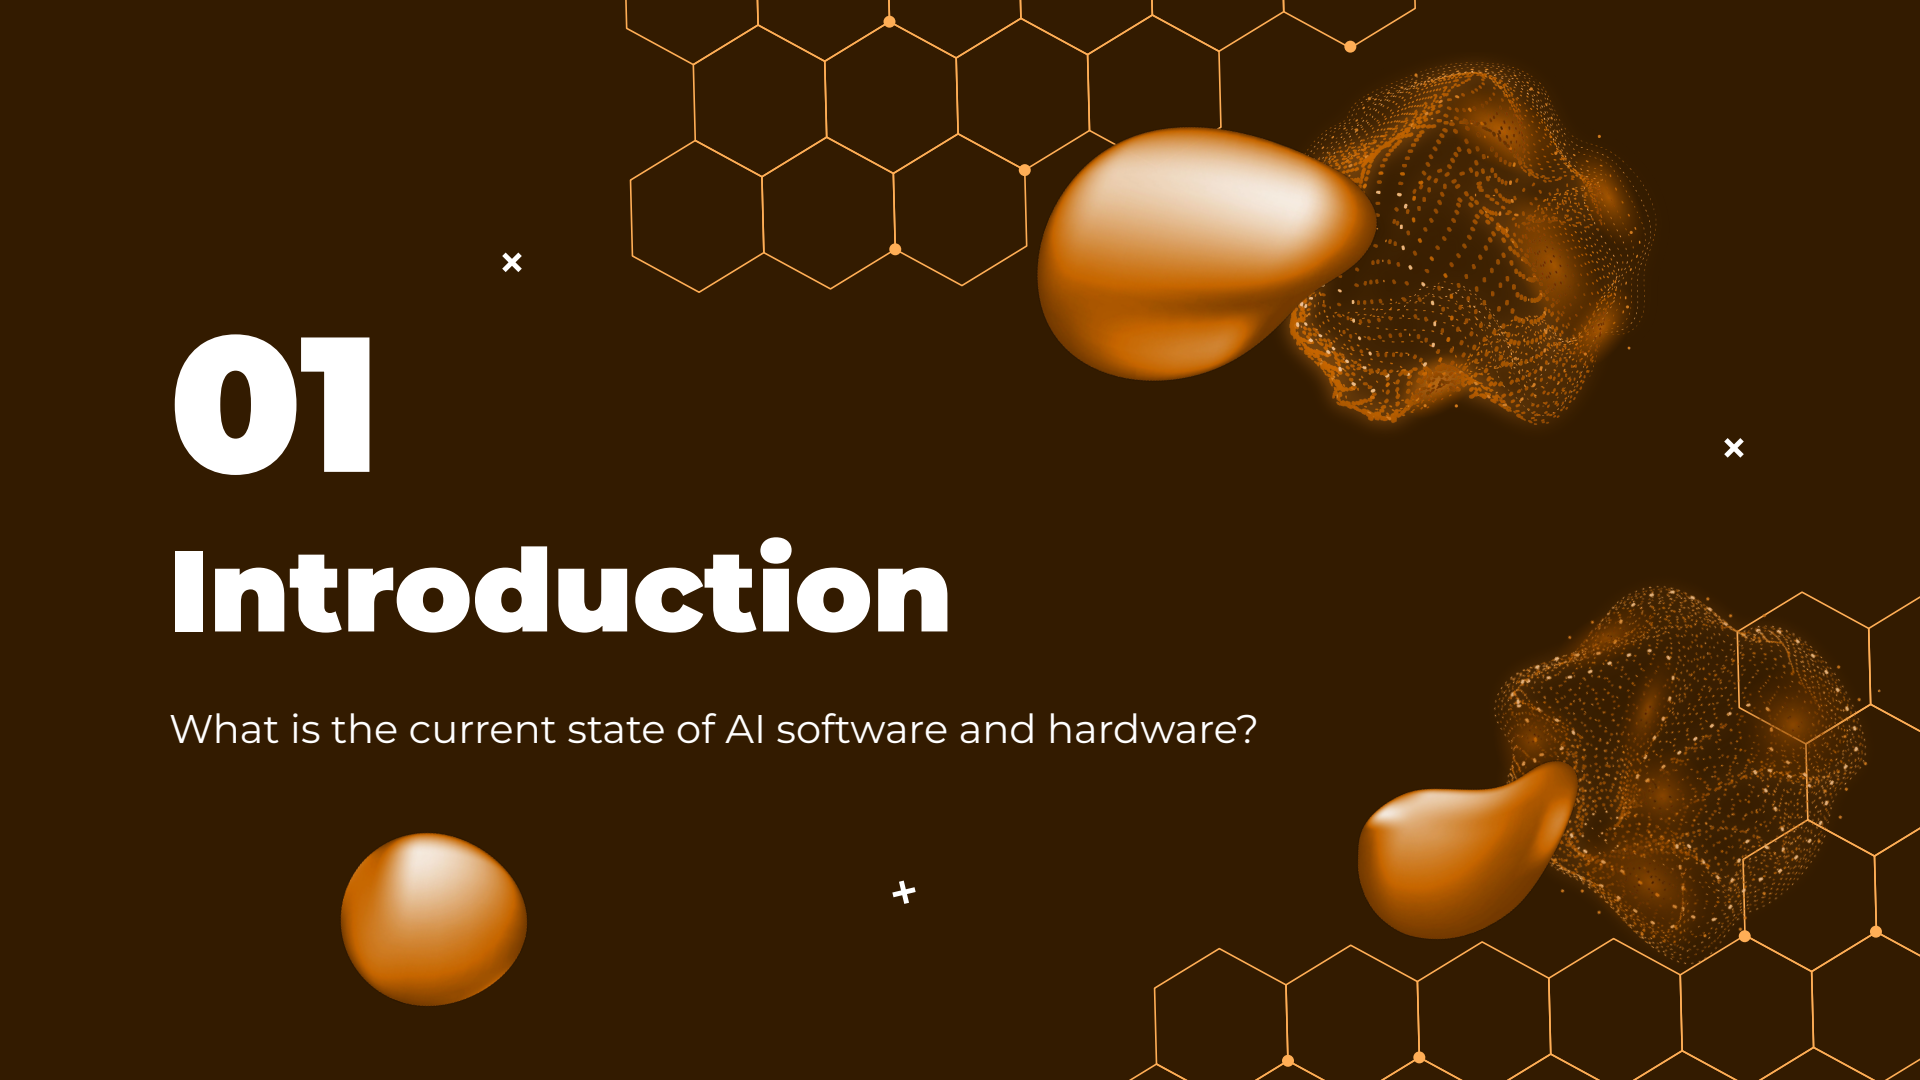

01 Introduction

What is the current state of AI software and hardware?

It’s a black box and/or hot glue

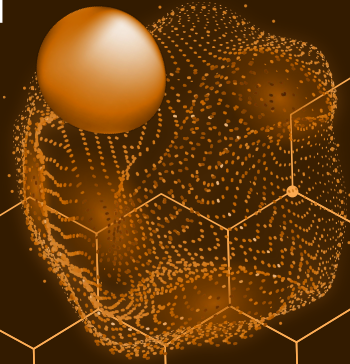

Apache TVM Path From Software To Hardware

Apache TVM: Compiler Framework

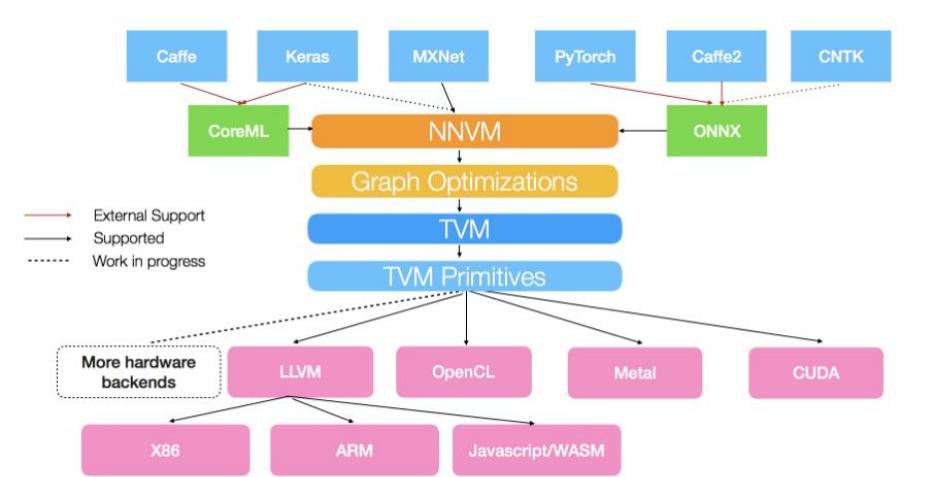

Fig. 1. What a modern ecosystem might look like. Apache TVM. (https://tvm.apache.org/2017/10/06/nnvm-compiler-announcement).

IREE: Path From Software To Hardware

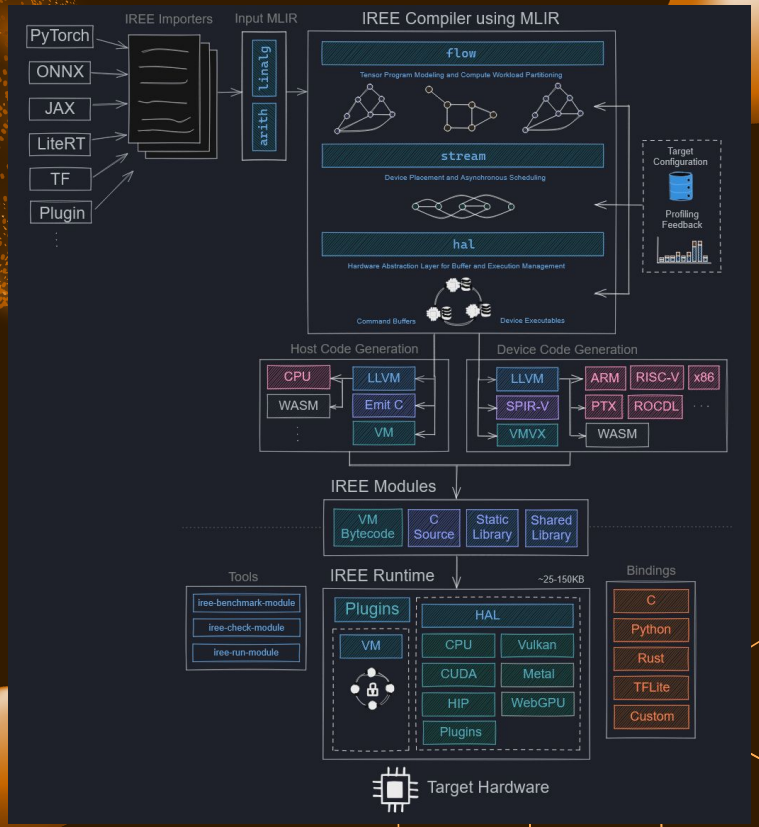

Two Language Problem

Python

High level, easy
SUPER SLOW
Calls into C++/ etc.

C++ / CUDA / PTX

Requires good
programming skills,
powerful, tedious

FLOPs go BRRR

New Hardware

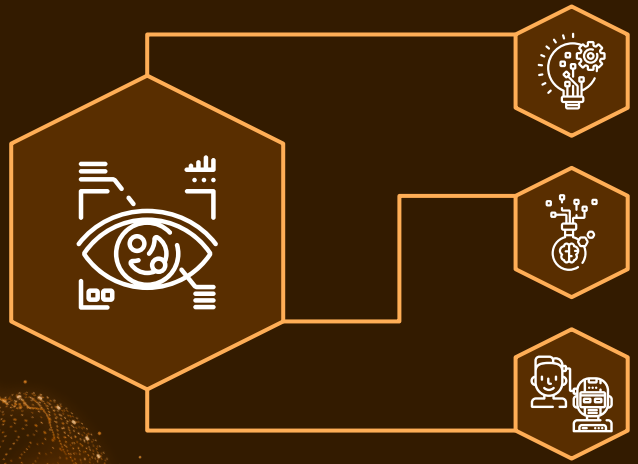

TPUs, LPUs — Google, Groq

Intel Loichi — SNNs, not tensors

Who knows? — We should future proof!

How has monopoly shaped our industry?

Lack of competition

Shared space

02 “The Hardware Lottery”

How has hardware shaped research?

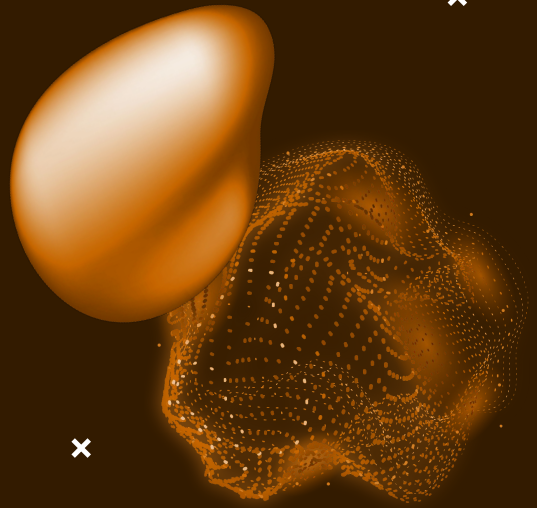

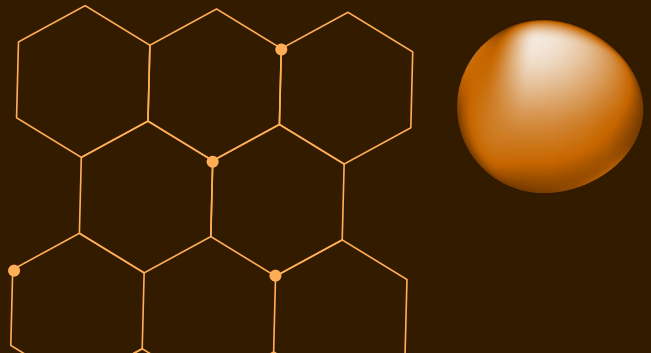

“This essay introduces the term hardware lottery to describe when a research idea wins because it is suited to the available software and hardware and not because the idea is superior to alternative research directions”

—Sara Hooker, Google Brain 2020

03 Current State

How do we deal with this now?

Compilers

IRs like MLIR, Auto-Tuning

[e]DSLs

Domain Specific Languages like Triton

Open Source

Design by committee failures, coordination failures, fragmentation

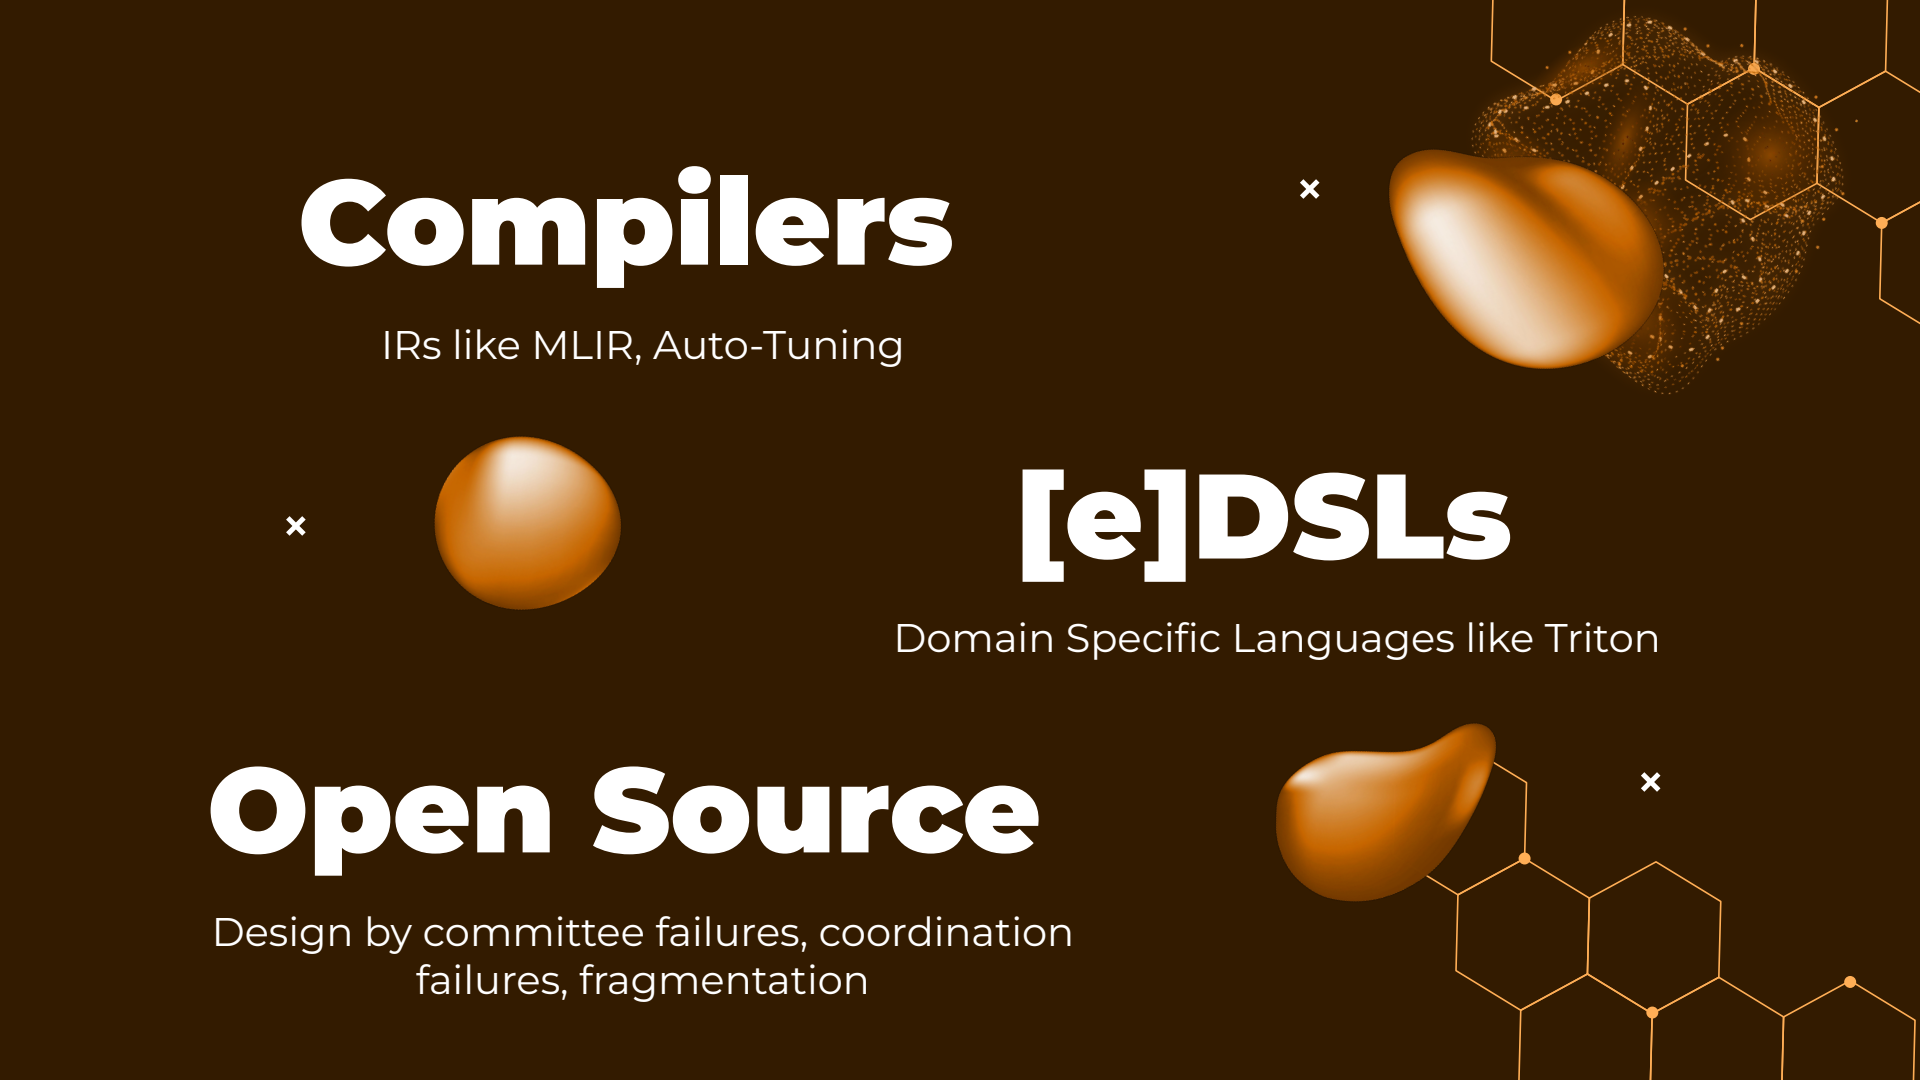

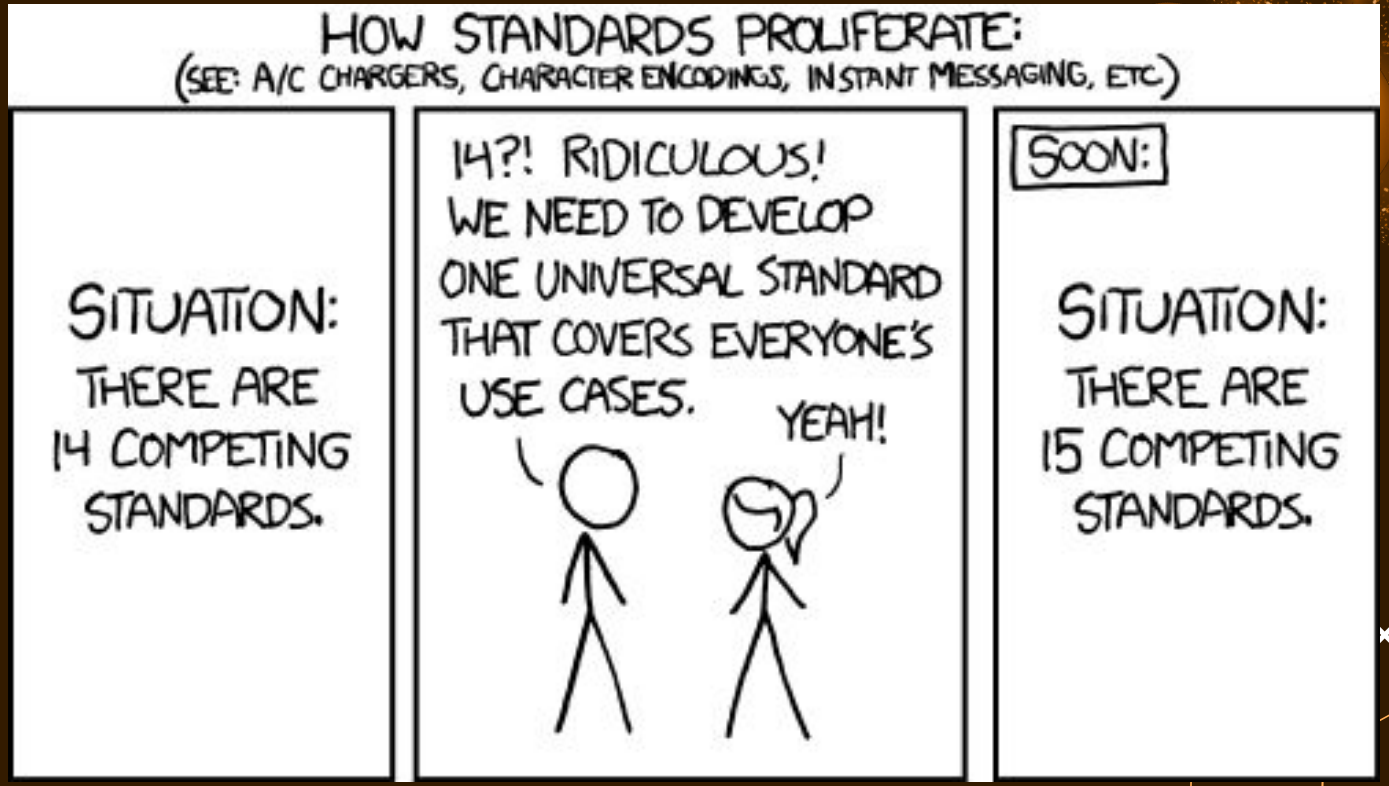
HOW STANDARDS PROLIFERATE:
(SEE: A/C CHARGERS, CHARACTER ENCODINGS, INSTANT MESSAGING, ETC)
SITUATION: THERE ARE 14 COMPETING STANDARDS.
14?! RIDICULOUS! WE NEED TO DEVELOP ONE UNIVERSAL STANDARD THAT COVERS EVERYONE'S USE CASES.
YEAH!
SOON:
SITUATION: THERE ARE 15 COMPETING STANDARDS.

This Means Vendor Lock-In

PTX for a single Nvidia GPU

OpenCL example - no Tensor Core support

Switching frameworks

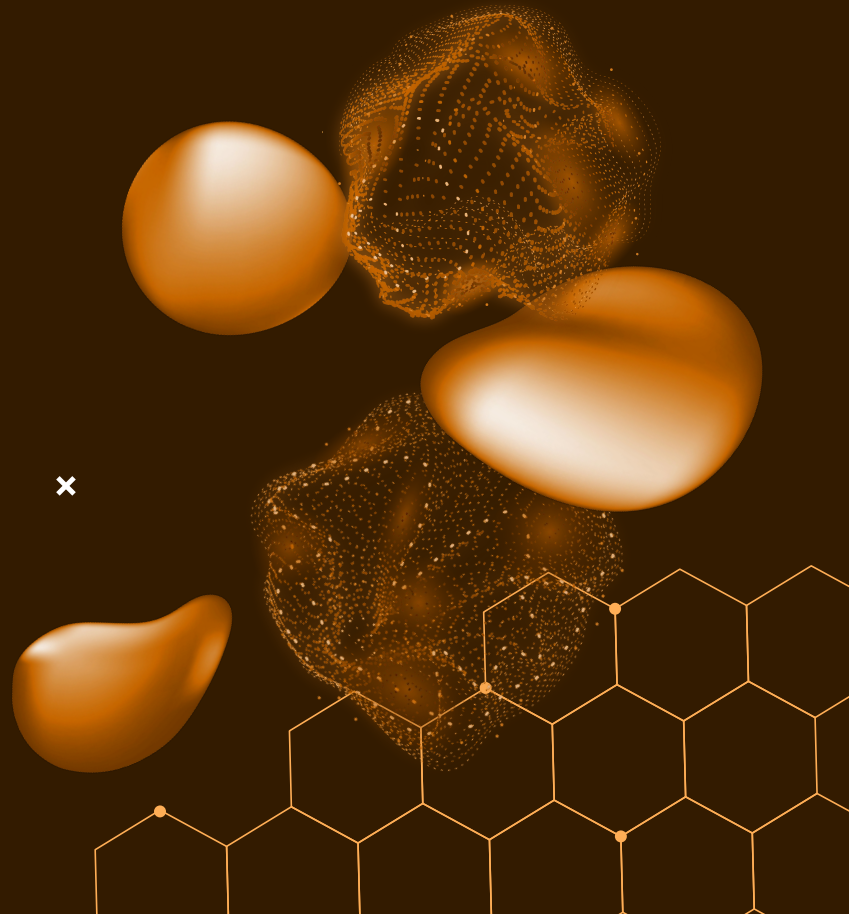


04
Modular / MAX / Mojo
Chris Lattner

Chris Lattner

2000	LLVM	IR, Clang
2010	Swift	Replaced Obj-C
2017	TF & TPU	Google

Compilers **Languages** **Hardware**

Modular

Scheduling, Batching, Cloud or Local Deployment

Max Serve

Graph compilers, libraries, tools

Max Engine

Python Superset - FAST

Mojo 🔥

Python Superset

No need to learn C++ / Rust / CUDA. Use Python libraries!

MLIR Coherency

Optimizations are easy across ALL stages

Typed, Compiled!

“Up to 65,000x faster than Python”

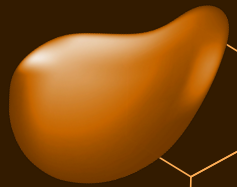

Mojo

Easy to write code! Performant!

Auto vectorizing
Auto parallelizing
Native SIMD
Easy tiling
GPU support just added
Only *slightly* verbose

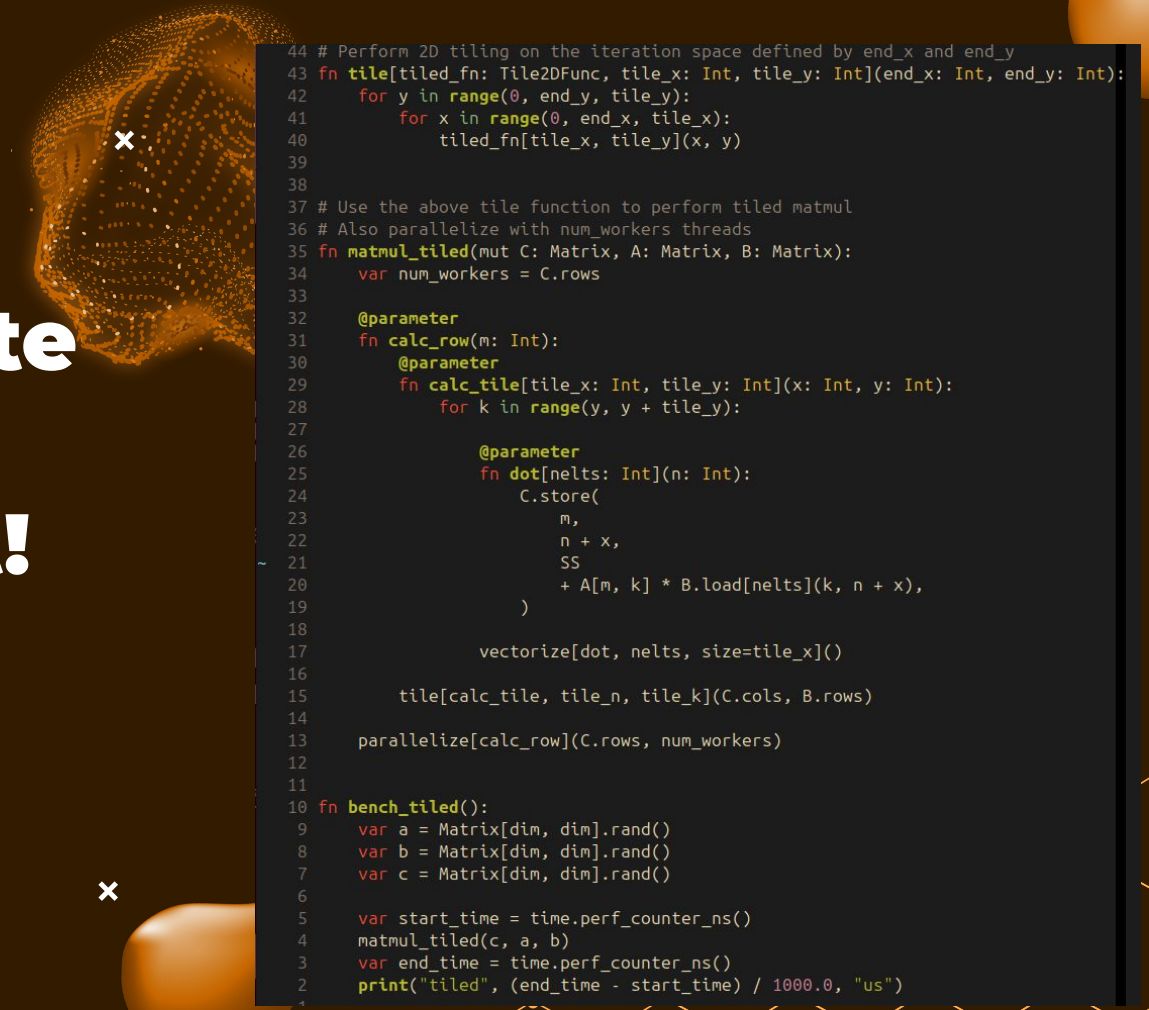

```
# Perform 2D tiling on the iteration space defined by end_x and end_y
fn tile[tiled_fn: Tile2DFunc, tile_x: Int, tile_y: Int](end_x: Int, end_y: Int):
    for y in range(0, end_y, tile_y):
        for x in range(0, end_x, tile_x):
            tiled_fn[tile_x, tile_y](x, y)


# Use the above tile function to perform tiled matmul
# Also parallelize with num_workers threads
fn matmul_tiled(mut C: Matrix, A: Matrix, B: Matrix):
    var num_workers = C.rows

    @parameter
    fn calc_row(m: Int):
        @parameter
        fn calc_tile[tile_x: Int, tile_y: Int](x: Int, y: Int):
            for k in range(y, y + tile_y):

                @parameter
                fn dot[nelts: Int](n: Int):
                    C.store(
                        m,
                        n + x,
                        SS
                        + A[m, k] * B.load[nelts](k, n + x),
                    )

                vectorize[dot, nelts, size=tile_x]()

        tile[calc_tile, tile_n, tile_k](C.cols, B.rows)

    parallelize[calc_row](C.rows, num_workers)


fn bench_tiled():
    var a = Matrix[dim, dim].rand()
    var b = Matrix[dim, dim].rand()
    var c = Matrix[dim, dim].rand()

    var start_time = time.perf_counter_ns()
    matmul_tiled(c, a, b)
    var end_time = time.perf_counter_ns()
    print("tiled", (end_time - start_time) / 1000.0, "us")
```

Matrix Multiplication 1024 x 1024

Language / Compiler	Notes	MicroSeconds
Python 3.12	Naive	50,800,000
Python 3.12	Numpy, Transpose	1,175,000
Mojo	Vectorized, Parallelized	8,300
C / Zig CC	Naive	1,950,000
C / Zig CC	-O3, OpenMP, Transpose	27,000

SOTA: MatMul & Relu

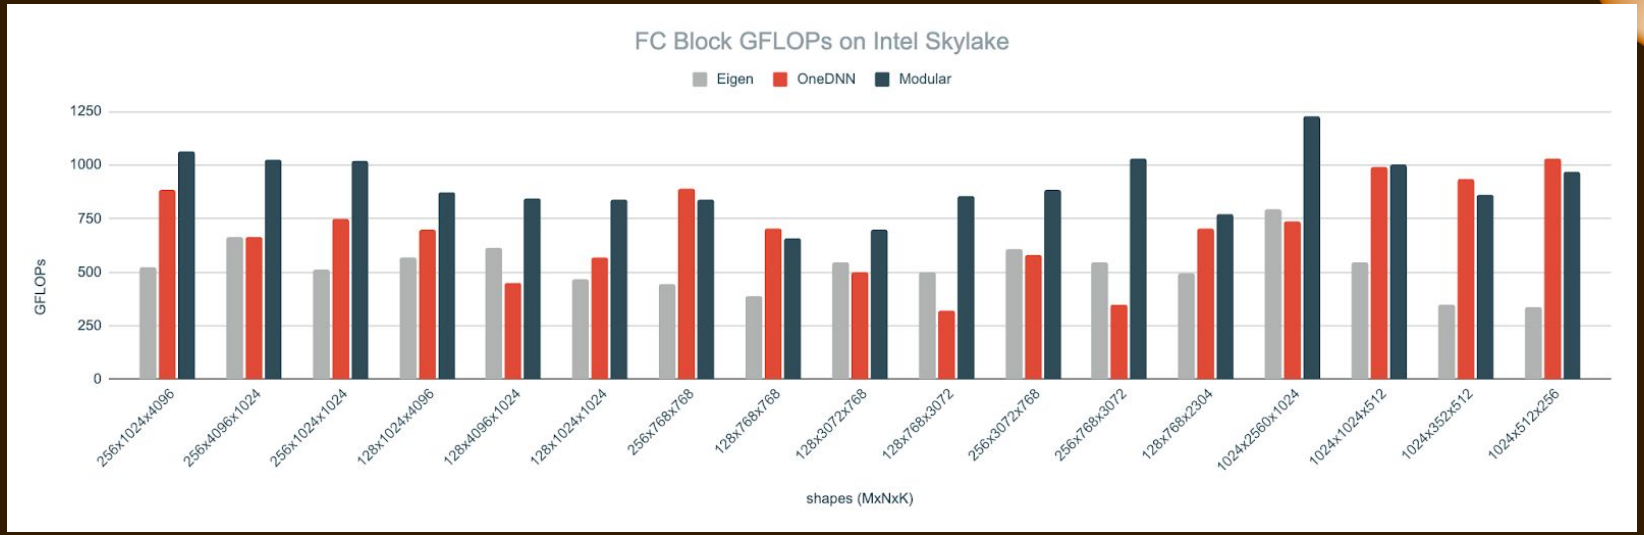

Critiques

“Ownership”

This memory management paradigm can be tough

New Keywords

Mojo is “Python ++” & WIP

Parameters

Separation of comptime is new to most

MAX System

Graph building here is a new system / library to learn

05

Takeaways

Chris Lattner-senapai notice me :3 UwU

Takeaways

Hardware

The Hardware Lottery

Ergonomics

Developer Experience

Unification

How do we UNITE research?

Future?

What's coming next?

THANKS!

DO YOU HAVE ANY QUESTIONS?

paul.dutton@my.utsa.edu

CREDITS: This presentation template was created by **Slidesgo**, and includes icons by **Flaticon**, and infographics & images by **Freepik**

References

Sara Hooker, “The Hardware Lottery”. 2020. https://arxiv.org/abs/2009.06489

https://tvm.apache.org/2017/10/06/nnvm-compiler-announcement

https://iree.dev/#project-architecture

https://xkcd.com/927/

Modular: The world's fastest unified matrix multiplication

https://static.wikia.nocookie.net/monopoly/images/4/41/Monopoly_2D_Art_Render.png/revision/latest?cb=20220109225628